

Distributed Systems - Lecture 4

Verteilte Systeme (BVS2)

Table of contents

- 2 Communication in Distributed Systems
 - 2.3 Programming with Sockets
- 3 Cooperation
 - 3.1 The Client-Server Model: Fundamental Properties
 - 3.2 Alternative: Peer-to-Peer Cooperation (P2P)

Zoom

<https://th-koeln.zoom-x.de/j/64426226965>

Meeting-ID: 644 2622 6965

Kenncode: 263566

Please post your questions or upvote existing questions during the lecture!

Source: Slides adapted from the BVS2 lecture with permission from Prof. Dr. Carsten Vogt

Summary (Last Week)

- **Message Passing**

- Direct or indirect (via mailbox or port)
- Synchronous (blocking) vs. asynchronous (non-blocking) communication
- “Rendezvous” = blocking send + blocking receive

- **Ports and Sockets**

- Ports: Logical endpoints for client-server coordination
- Sockets: API to protocol stacks (TCP/IP, UDP/IP)
 - Domains: UNIX (local), Internet (network)
 - Types: Stream (TCP), Datagram (UDP), Raw (manual)

- **Client-Server Communication**

- Server socket: Waits for connections
- Client socket: Initiates connection or sends datagrams

- **Programming Model**

- Setup: Bind → Listen → Accept (server); Connect (client)
- Supports both local (UNIX) and network (Internet) communications


2 Communication in Distributed Systems

2.3 Programming with Sockets

Stream Sockets: Program Steps

1. Creation of server sockets:

`socket()`, `bind()`, `listen()`



2. Creation of client sockets:

`socket()`



3. Request connection

`connect()`


4. Connection:

`accept()`

5. Data connection: `sendall()`/`recv()`

6. Closing the socket: `close()`

Stream Sockets - Timeline



Internet Domain Stream Sockets Example: Client

```
1 import socket
2
3 HOST = 'nebsy.nt.fh-koeln.de' # Server's IP address (use localhost for local testing)
4 PORT = 62423                 # Server's port
5
6 def send_message(message):
7     with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as client_socket:
8         client_socket.connect((HOST, PORT))
9         client_socket.sendall(message.encode())
10        response = client_socket.recv(1024)
11        print("Server response:", response.decode())
12
13 if __name__ == "__main__":
14     msg = input("Enter a message to send to the server: ")
15     send_message(msg)
```

Internet Domain Stream Sockets Example: Server

```
1 import socket
2
3 HOST = '0.0.0.0' # Listen on all available interfaces
4 PORT = 62423      # Port to listen on
5
6 def start_server():
7     with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as server_socket:
8         server_socket.bind((HOST, PORT))
9         server_socket.listen()
10        print(f"Server listening on port {PORT}...")
11
12    while True:
13        conn, addr = server_socket.accept()
14        with conn:
15            print(f"Connected by {addr}")
16            while True:
17                data = conn.recv(1024)
18                if not data:
19                    break
20                response = f"Received message: {data.decode()}"
21                conn.sendall(response.encode())
22
23    ..
```


(See uploaded Python files for full code listing)

Stream Sockets - Programming Remarks

- Server can send a result back to the client over the same connection
- Server: `conn.sendall(response.encode())`
- Client: `client_socket.recv(1024)`
- If the number of bytes is unknown:

```
1 # Buffer to hold the received data
2 received_data = b""
3
4 # Keep receiving data until the server closes the connection
5 while True:
6     chunk = sock.recv(4096) # You can adjust buffer size as needed
7     if not chunk:
8         # No more data (connection closed)
9         break
10    received_data += chunk
11
12 print("Received data:")
13 print(received_data.decode('utf-8', errors='replace'))
```

Datagram Sockets - Timeline



Internet Domain Datagram Sockets Example: Client

```
1 import socket
2
3 def udp_client(server_host='127.0.0.1', server_port=12345):
4     with socket.socket(socket.AF_INET, socket.SOCK_DGRAM) as client_socket:
5         while True:
6             message = input("Enter message (or 'exit' to quit): ")
7             if message.lower() == 'exit':
8                 break
9
10            client_socket.sendto(message.encode(), (server_host, server_port))
11            data, _ = client_socket.recvfrom(1024)
12            print(f"Received from server: {data.decode()}")
13
14 if __name__ == "__main__":
15     udp_client()
```

Internet Domain Datagram Sockets Example: Server

```
1 import socket
2
3 def udp_server(host='127.0.0.1', port=12345):
4     with socket.socket(socket.AF_INET, socket.SOCK_DGRAM) as server_socket:
5         server_socket.bind((host, port))
6         print(f"UDP server listening on {host}:{port}")
7
8     while True:
9         data, addr = server_socket.recvfrom(1024) # Buffer size is 1024 bytes
10        print(f"Received from {addr}: {data.decode()}")
11
12        response = f"Echo: {data.decode()}"
13        server_socket.sendto(response.encode(), addr)
14
15 if __name__ == "__main__":
16     udp_server()
```

(See uploaded Python files for full code listing)

Raw Sockets

The application must create the packet header itself and can read the packet header

- This is in contrast to stream sockets and datagram sockets where this is automatically handled by the protocol implementation
- Advantage: The application has access to the headers
- e.g., for experimenting with header entries
- e.g., for sniffing

Select and Poll

- A single functional call that waits for incoming data on several sockets
- Enables server to serve several clients
- Can add timeouts

Select

- `select.select()` monitors multiple sockets (or files) for readability, writability and errors
- Useful in non-blocking or asynchronous I/O

```
1 import select
2 readable, writable, errored = select.select(inputs, outputs, inputs, timeout)
```

Poll

- `poll()` is more scalable than `select()` (better for many sockets)
- Works with a poll object and file descriptors

```
1 poller = select.poll()
2 poller.register(socket.fileno(), select.POLLIN)
```


- Events are checked using `poller.poll(timeout)`

Select Example

```
1 import socket, select
2
3 server = socket.socket()
4 server.bind(('localhost', 12345))
5 server.listen()
6
7 inputs = [server]
8
9 while True:
10     readable, _, _ = select.select(inputs, [], [])
11     for s in readable:
12         if s is server:
13             conn, addr = server.accept()
14             inputs.append(conn)
15         else:
16             data = s.recv(1024)
17             if data:
18                 print("Received:", data.decode())
19             else:
20                 inputs.remove(s)
21                 s.close()
```

3 Cooperation

Communication and Cooperation



- Often: cooperation based on the **client-server** paradigm
 - In this chapter:
 - Fundamental properties of client-server systems
 - Briefly: **peer-to-peer** as an alternative
 - Implementation techniques for client-server systems

3.1 The Client-Server Model: Fundamental Properties

Client-Server System: Roles


Roles for Processes / Threads / Computer Nodes:

- **Server**: Provider of services
 - e.g., access to databases, file systems, documents
- **Client**: User of services
 - Calls services locally or over the network

Dual Role:


- A process / thread / computer node can be both:
 - A **server**: provides some services
 - And a **client**: uses other services

Client-Server Cooperation: Steps




Client-Server Cooperation: Timing Behaviour

a.) Cooperation without reply:




b.1.) Synchronous cooperation with reply




Client-Server Cooperation: Timing Behaviour

b.1.1.) Asynchronous cooperation with reply: Client polling for the reply



b.1.2.) Asynchronous cooperation with reply: Reply interrupting the client



Client-Server Cooperation: Possible Errors

- Messages can be:
 - Lost
 - Distorted
 - Duplicated
- The server might be:
 - Down
 - Unreachable

Consequences:

- It is **not guaranteed** that a request:
 - Is processed at all
 - Is not processed **multiple times**
- Additional mechanisms are needed to provide different types of **guarantees**

Client-Server Cooperation: Guarantees & Mechanisms

“At Most Once” Guarantee

- A request is processed **at most once**
- **Mechanism:** Sequence numbers
 - Each request has a unique number
 - Server discards requests with numbers it has already seen

“At Least Once” Guarantee


- A request is processed **at least once**
- **Mechanism:** Acknowledgments
 - Server acknowledges receipt of a request
 - Client resends requests not acknowledged within a timeout
 - May resend to a different server

“Exactly Once” Guarantee

- A request is processed **exactly once**
- **Mechanism:** Combination of the above two

3.2 Alternative: Peer-to-Peer Cooperation (P2P)

- Files / services are distributed over „peer nodes“
- Each node can use services of the other nodes



Alternative: Peer-to-Peer Cooperation (P2P)

P2P: Decentralized Organization

- Client-server: centralized organization with a few servers
- P2P: decentralized, all participants can act as both client and server

Advantages

- Many resources available
- Load distribution and fault tolerance
 - Achieved by replicating resources
- Self-organization

Drawbacks

- Higher costs for searches and communication
- Security and data protection are harder to enforce

Summary

- **Stream Sockets (TCP):**

- Connection-oriented: server listens and accepts, client connects.
- Communication via `sendall()` and `recv()`.

- **Datagram Sockets (UDP):**

- Connectionless: each `sendto()/recvfrom()` is a discrete message.
- No guaranteed delivery/order—requires manual error handling.

- **Client-Server Model:**

- Server = service provider; Client = service consumer.
- Cooperation sequence: request → processing → response.

- **Communication Guarantees:**

- *At most once*: sequence numbers. *At least once*: acknowledgments + retries.
- *Exactly once*: combine both mechanisms.

- **P2P Alternative:**

- Fully decentralized, each node is both a client and a server.
- Benefits: redundancy, scalability. Challenges: discovery, security, consistency.

In two weeks

- More on the client-server model
- Remote Procedure Call